




The bug was in `model.compile()` in Section 6. The loss function was set to 'mse' and the metric was set to 'mae', which are mainly used for regression. Since this is a multi-class classification problem with one-hot encoded labels and a softmax output layer, I changed the loss to 'categorical_crossentropy' and the metric to 'accuracy'. I also updated METRIC_TO_PLOT in Section 7 from 'mae' to 'accuracy' to keep the training plots consistent with the changes I made.